

Copula Estimation - Univariate Case

2025-11-27

1. Sampling with copula package

The copula package automatically samples bivariate and multivariate copulas with a convenient function. To use it in our univariate scenario, we consider the number of observations in the dataset as the dimension of the copula. Below the Gumbel case:

```
set.seed(46)
library(copula)
theta <- 1.5
n <- 200

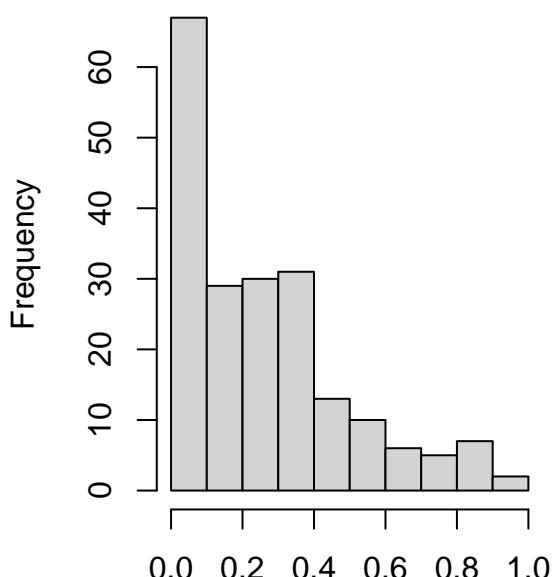
Gcop <- gumbelCopula(param = theta, dim = n)

U_cop <- rCopula(1, Gcop)
```

Choosing the margin allows to construct a sample $X_1, \dots, X_n \sim H = C(F(x_1), \dots, F(x_n))$ by applying $X_i = F^{-1}(U_i)$. For example, taking $X_i \sim Frechet(2)$:

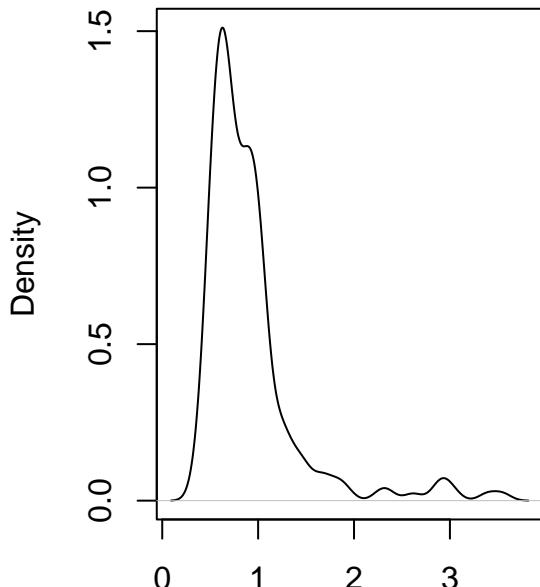
```
alpha <- 2
X_cop <- qfrechet(U_cop, shape = alpha)
```

Gumbel – copula package



U_{cop}

Density of Gumbel–Frechet



$N = 200$ Bandwidth = 0.09007

2. Stochastic Representation with latent V

To sample from the same copula, we can use the latent variable representation. The latent variable V whose Laplace transform is the Gumbel generator $\psi(t) = \exp\{-t^{\frac{1}{\theta}}\}$ has distribution $F_V \sim \text{Stable}(\alpha = 1/\theta, \beta = 1, \gamma = (\cos(\frac{\pi}{2\theta}))^\theta, \delta = 0; pm = 1)$.

```
set.seed(46)
library(stabledist)

gum_psi <- function(t, theta){
  exp(-t ^ (1/theta))
}

## Stable parameters for gumbel

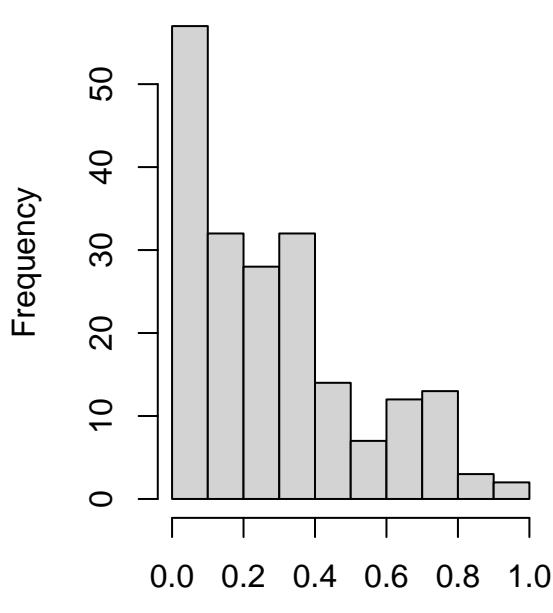
V <- rstable(
  n = 1,
  alpha = 1/theta,
  beta = 1,
  gamma = cospi(1/(2 * theta)) ^ theta,
  delta = 0,
  pm = 1
)

E <- rexp(n, V)

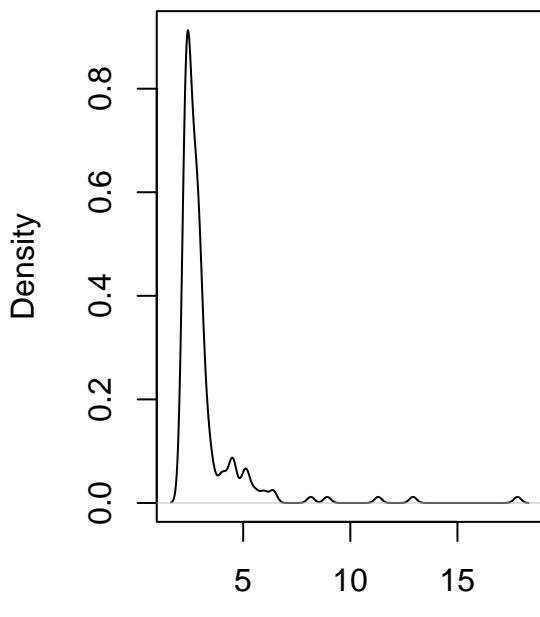
U_v <- gum_psi(E, theta)

X_v <- qfrechet(U_v, alpha)
```

Gumbel – Latent Variable



Density of Gumbel–Frechet



N = 200 Bandwidth = 0.1699

For convenience, we write the sampler function for the U values (the quantile function is left unspecified to allow the freedom to choose the margin F):

```
# Latent variable
rGumbV <- function(n, theta) {
  require(stabledist)

  # Gumbel V r.v.
  V <- rstable(
    n = 1,
    alpha = 1/theta,
    beta = 1,
    gamma = cospi(1/(2 * theta))^theta,
    delta = 0,
    pm = 1
  )

  E <- rexp(n, V)

  # Gumbel generator
  gum_psi <- function(t, theta){
    exp(- t ^ (1/theta))
  }

  #(U_1, ..., U_n) ~ C_psi
  U <- gum_psi(E, theta)

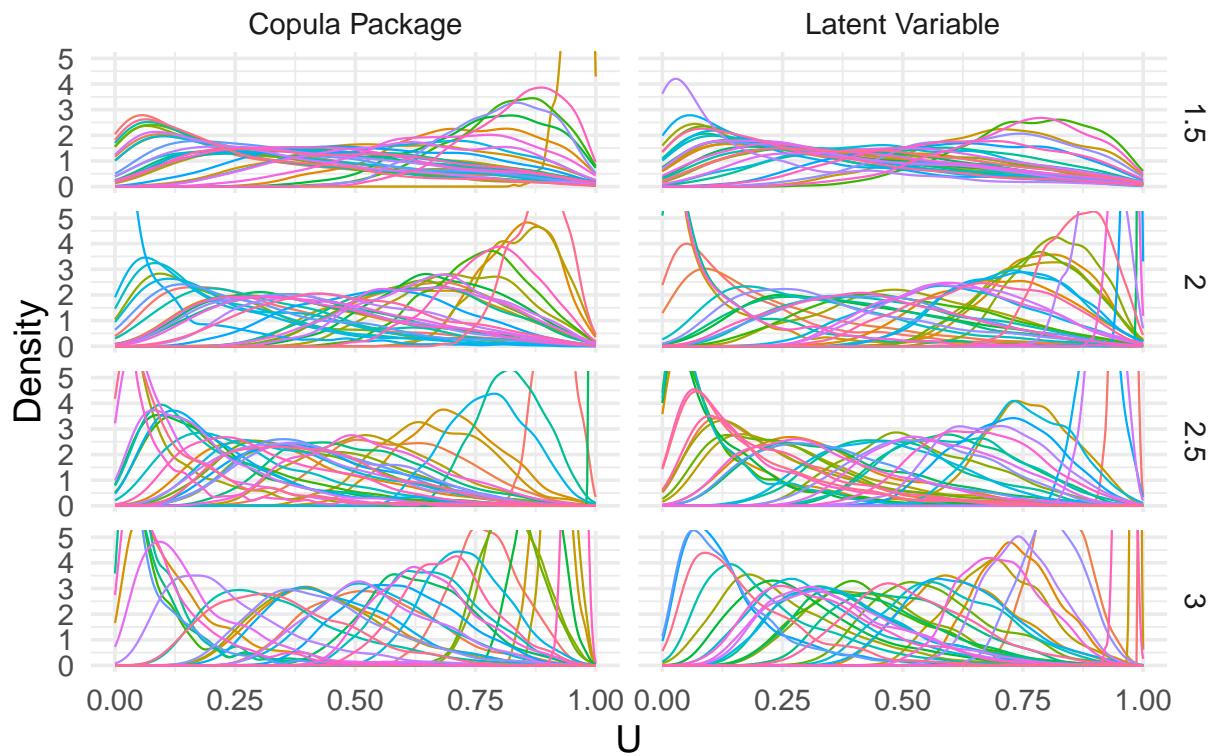
  return(U)
}
```

3. Model likelihood and identifiability

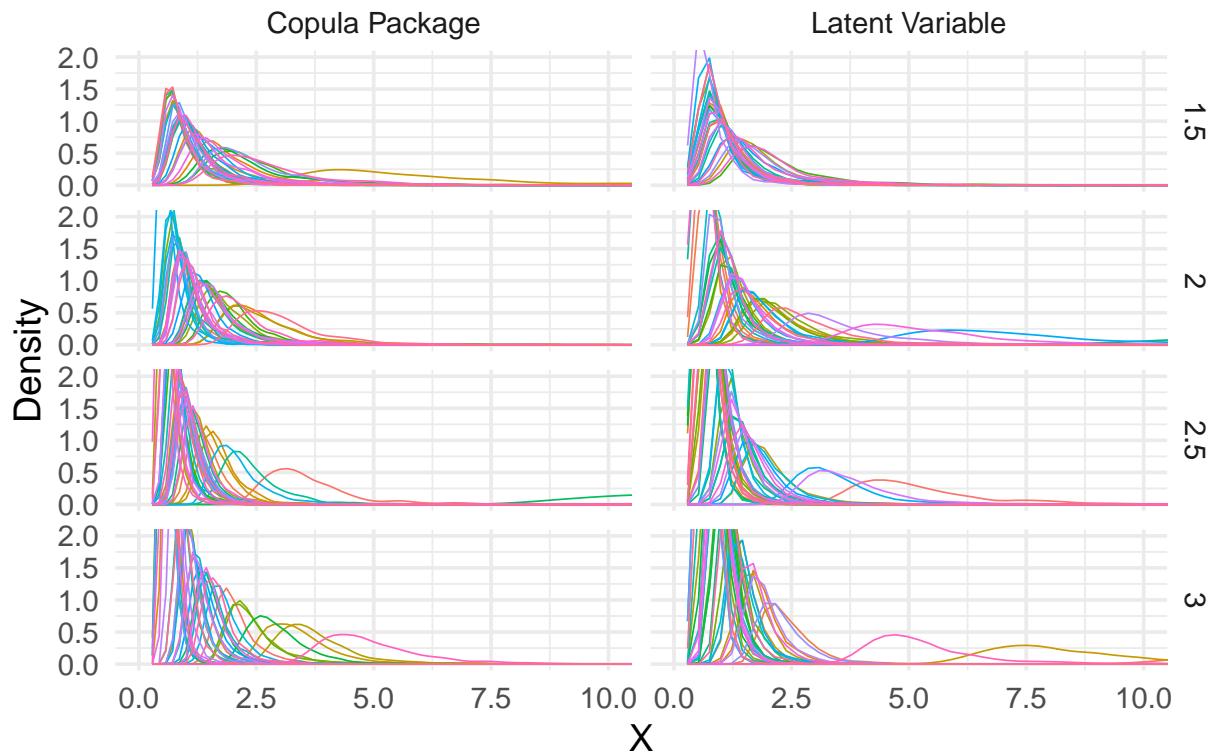
3.1 Simulation

Below we can see different realisations from a Gumbel-Frechet model and different theta parameters. It seems that different theta parameters lead to very similar realisations. When we perform estimation of the parameter, could this create non identifiability?

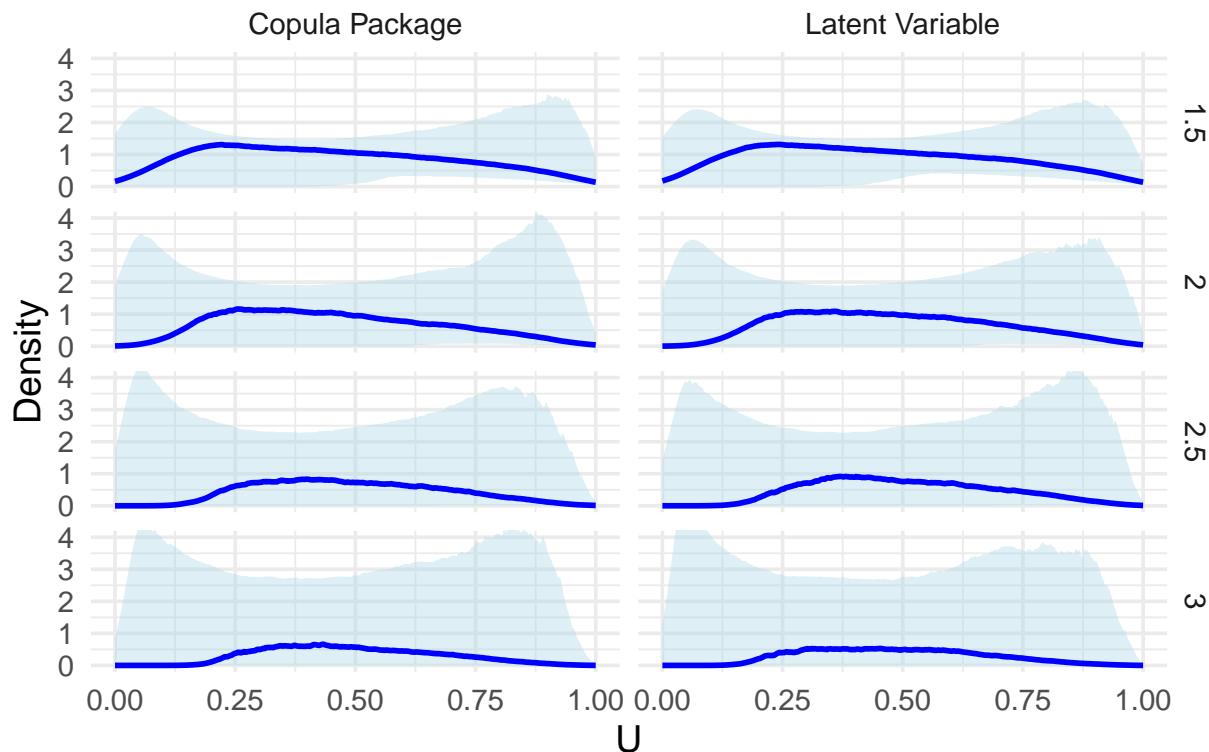
Gumbel – 30 replicates



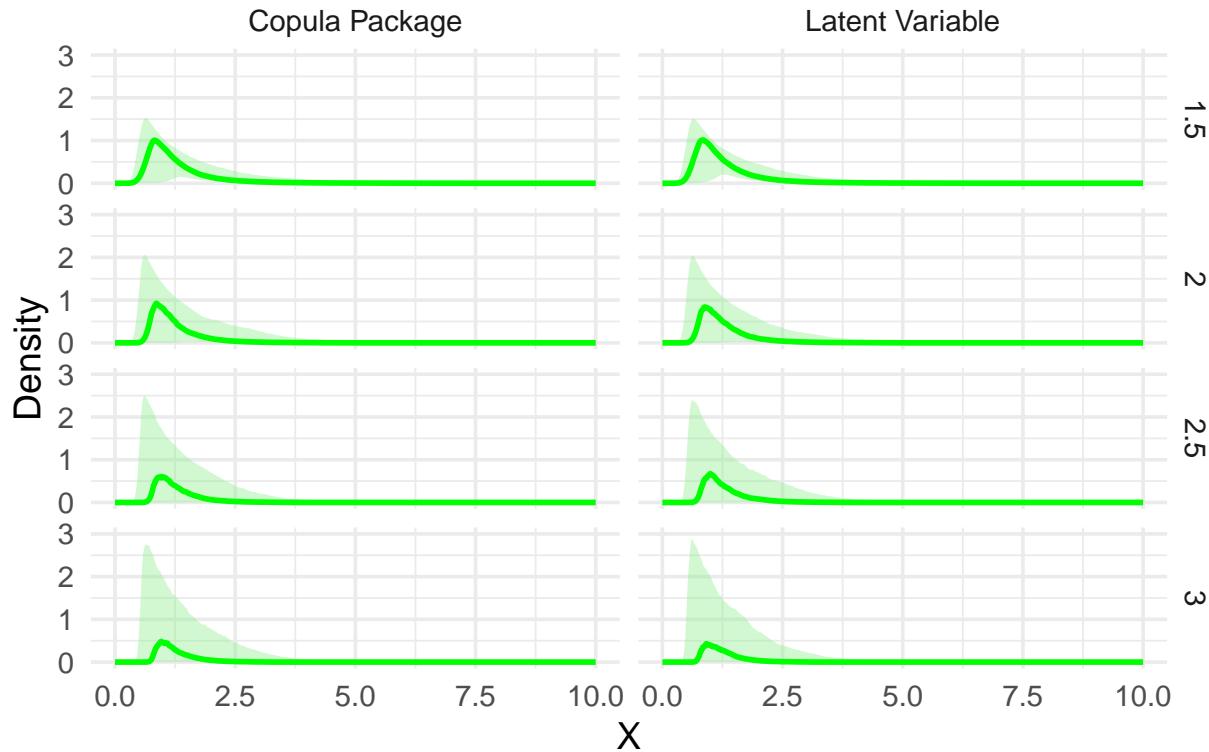
GFrechet – 30 replicates



Gumbel – 1000 replicates – Median and 80%



GFrechet – 1000 replicates – Median and 80%



3.2 Likelihood

The aim is to show what is the shape of the likelihood of the model.

$$\text{Model: } (X_1, \dots, X_n) \sim H_n(\cdot) \quad \text{with} \quad H_n(x_1, \dots, x_n) = C_\theta(F(x_1), \dots, F(x_n)),$$

where F is the marginal CDF, f its density, and C_θ a copula (with copula density c_θ).

$$\begin{aligned} H_n(x_1, \dots, x_n) &= C_\theta(u_1, \dots, u_n), \quad u_i := F(x_i) \\ \implies h_n(x_1, \dots, x_n) &= \frac{\partial^n}{\partial x_1 \cdots \partial x_n} H_n(x_1, \dots, x_n) \\ (\text{chain rule}) \implies &\frac{\partial^n}{\partial u_1 \cdots \partial u_n} C_\theta(u_1, \dots, u_n) \cdot \prod_{i=1}^n \frac{\partial u_i}{\partial x_i} \\ \Rightarrow h_n(x_1, \dots, x_n) &= c_\theta(F(x_1), \dots, F(x_n)) \prod_{i=1}^n f(x_i) \end{aligned} \quad (1)$$

Now the likelihood (viewed as a function of the copula parameter θ) for the observed vector (x_1, \dots, x_n) is

$$L(\theta) = h_n(x_1, \dots, x_n; \theta) = c_\theta(F(x_1), \dots, F(x_n)) \prod_{i=1}^n f(x_i). \quad (2)$$

Taking logarithms yields the log-likelihood

$$\ell(\theta) := \log L(\theta) = \log c_\theta(F(x_1), \dots, F(x_n)) + \sum_{i=1}^n \log f(x_i). \quad (3)$$

Note that in the case of estimating only θ , the log-likelihood is constant in the second term w.r.t. the parameter. Note as well that for a too large sample size, the evaluation of the copula density is not feasible.

```
cop <- gumbelCopula(param = theta_true, dim = 1000)
U <- as.numeric(rCopula(1, cop))
dCopula(U, copula = cop)
```

```
## Error in polyG(lx[i.m[[meth]]], alpha = alpha, d = d, method = meth, log = log): d > 220 not yet supp
```

To illustrate what happens to the MLE, we can compute the profile log-likelihood for a given sample with the following implementation. A crucial factor seems to be the estimation of the margin to obtain the U-values used to estimate the parameter from the log-likelihood.

```
set.seed(123)

# -----
# 1. TRUE parameter and dimension
# -----
n <- 100
theta_true <- 3

# -----
# 2. Simulate ONE 100-dimensional vector from Gumbel copula
# -----
cop <- gumbelCopula(param = theta_true, dim = n)
```

```

U <- as.numeric(rCopula(1, cop))    # true uniforms

# -----
# 3. Generate X via lognormal margins
# -----
mu <- 0
sigma <- 1
X <- qlnorm(U, meanlog = mu, sdlog = sigma)

# -----
# 4A. Parametric margins
# -----
mu_hat <- mean(log(X))
sigma_hat <- sd(log(X))
u_hat_param <- plnorm(X, meanlog = mu_hat, sdlog = sigma_hat)

# -----
# 4B. ECDF pseudo-observations
# -----
u_hat_ecdf <- rank(X) / (n + 1)

# -----
# 5. Gumbel log-likelihood
# -----
loglik_gumbel <- function(theta, u) {
  if (theta <= 1) return(-1e10)  # keep optimizer stable
  cop <- gumbelCopula(param = theta, dim = length(u))
  ll <- log(dCopula(u, copula = cop))
  return(ll)
}

# Negative log-likelihood for optimization
negLL <- function(theta, u) -loglik_gumbel(theta, u)

# -----
# 6. Compute estimators via optim()
# -----
# TRUE UNIFORMS
est_trueU <- optim(par = 5, fn = negLL, u = U, method="L-BFGS-B",
                     lower = 1.001, upper = 30)$par

# PARAMETRIC MARGINS
est_param <- optim(par = 5, fn = negLL, u = u_hat_param, method="L-BFGS-B",
                     lower = 1.001, upper = 30)$par

# ECDF PSEUDO-OBSERVATIONS
est_ecdf <- optim(par = 5, fn = negLL, u = u_hat_ecdf, method="L-BFGS-B",
                     lower = 1.001, upper = 30)$par

# -----
# 7. Likelihood curves for comparison
# -----

```

```

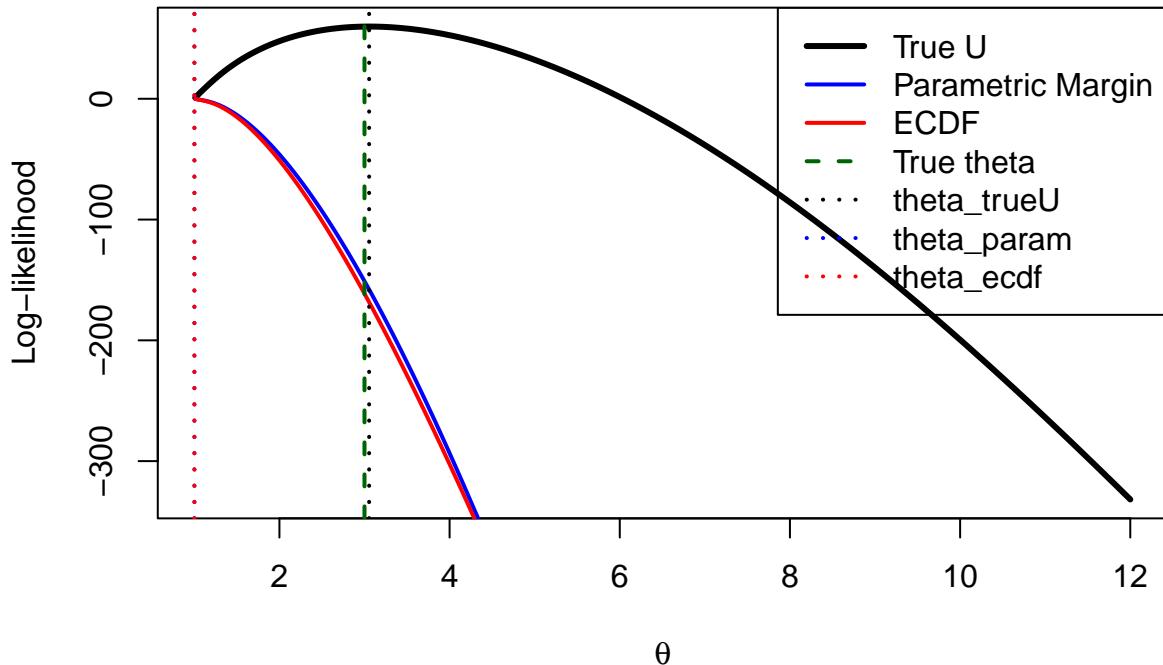
theta_grid <- seq(1.01, 12, length.out = 200)

ll_trueU <- sapply(theta_grid, loglik_gumbel, u = U)

## Loading required namespace: Rmpfr
ll_param <- sapply(theta_grid, loglik_gumbel, u = u_hat_param)
ll_ecdf <- sapply(theta_grid, loglik_gumbel, u = u_hat_ecdf)

```

Gumbel log-likelihood shapes



```

##  

## ESTIMATED THETAS:  

## theta_true = 3  

## theta_trueU    = 3.054531  

## theta_param   = 1.001  

## theta_ecdf    = 1.001

```